

Education series - August 2026

What ~100 go-to-market leaders told us about AI agents *and the data behind it*

Six patterns from conversations with go-to-market leaders across RevOps, sales, and marketing at mid-market and enterprise teams, each paired with the public research that says they're structural, not anecdotal.

~100

Conversations with go-to-market leaders

6

Patterns that surfaced again and again

Executive summary

Six things we kept hearing from go-to-market leaders, *in their own words*

Qualitative themes, honest framing. These are patterns from ~100 conversations with go-to-market leaders, not a quantified survey, so we're not reporting percentages from the interviews. What we *are* reporting is what we heard, alongside third-party public data that corroborates it. Every statistic in this deck is external research, attributed to its source.

01 "Everyone's using it. No one can tie it to the number."

02 "The AI is only as good as our data, and our data's a mess."

03 "Every tool has its own AI, and none share a brain."

04 "I can't put an agent in front of a customer without guardrails."

05 "It never compounds. Every quarter starts from zero."

06 "The pilot worked. Production didn't."

Theme 01 · Attribution

"Everyone's using it. No one can tie it to the number."

Most pilots never defined success, so teams measure activity, not revenue. There's no view that ties an agent's work to a deal stage.

3 in 10

Agent teams aren't evaluating their agents at all

Only 52.4% run offline evals before deployment; just 37.3% evaluate live traffic. Most agent failures aren't mysterious; they're unmeasured.

LANGCHAIN, STATE OF AGENT ENGINEERING, DECEMBER 2025

39%

Of organizations see any EBIT impact from AI, though 88% use it

Most attribute less than 5% of EBIT to AI. Revenue gains show up most often in marketing and sales, when anyone can measure them.

MCKINSEY, THE STATE OF AI, NOVEMBER 2025

"The AI is only as good as our data, and our data's a mess."

The number-one barrier leaders name isn't the model. It's stale, scattered, ungrounded data underneath it.

60%

Of AI projects will be abandoned without AI-ready data

63% of data leaders either don't have, or aren't sure they have, the right data practices for AI.

GARTNER, FEBRUARY 2025

16% → 54%

Accuracy jump from knowledge-graph grounding

Same model, better foundation.
Structure is the accuracy mechanism.

DATA.WORLD AI LAB, ARXIV, 2023

84%

Of data leaders say their data strategy needs a complete overhaul before their AI ambitions can succeed

And 89% of leaders with AI in production have already seen inaccurate or misleading outputs.

SALESFORCE, STATE OF DATA & ANALYTICS, NOVEMBER 2025

Theme 03 · Fragmentation

"Every tool has its own AI, and none share a brain."

Per-agent memory means five tools give five answers and re-learn the same context. AI is being used, but rarely orchestrated.

66.4%

Of agentic AI is already multi-agent

Most teams aren't running one agent; they're running many. The shared-brain problem is the default condition, not the edge case.

LANDBASE / MARKET.US, 2025

100K → 97M

MCP installs in 16 months

The connective plumbing exploded: MCP is now under the Linux Foundation with OpenAI, Google, Microsoft, and AWS behind it. But shared context didn't come with the pipes.

ANTHROPIC, DECEMBER 2025

Theme 04 · Trust

"I can't put an agent in front of a customer without guardrails."

Legal and the CRO say no until there's role-based access, human-in-the-loop, and an audit trail. Until then, agents stay stuck in read-only demos.

1 in 4

Leaders completely trust their AI systems, yet two-thirds have given AI write access to core systems

Trust is trailing autonomy. The market is voting for governed autonomy: access with rules, not access instead of them.

KYNDRYL PEOPLE READINESS REPORT, JUNE 2026

22% → 63%

Leaders requiring human validation of agent outputs nearly tripled in a year

Human-in-the-loop went from afterthought to default in four quarters.

KPMG AI QUARTERLY PULSE, MARCH 2026

Theme 05 · Compounding

"It never compounds. Every quarter starts from zero."

Winners learn from feedback and retain context. Everyone else redoes the pilot, and the corrections never stick.

95%

Of enterprise GenAI pilots return zero, and the core barrier is learning

MIT's diagnosis: most systems "do not retain feedback, adapt to context, or improve over time." 66% of executives want AI that learns from feedback; 63% demand tools that retain context.

MIT PROJECT NANDA, THE GENAI DIVIDE, 2025

16%

Of enterprise AI deployments qualify as true agents

Menlo's bar: systems that plan, act, observe feedback, and adapt. The rest need a foundation that learns before they can.

MENLO VENTURES, DECEMBER 2025

Theme 06 · Production

"The pilot worked. Production didn't."

The demo impressed everyone. Then it met real data, real governance, and real workflows.

95%

Of AI agents never make it to production

The difference between the 95% and the 5% is rarely the model; it's the operating layer underneath it.

**CAPGEMINI RESEARCH INSTITUTE,
APRIL 2026**

2%

Of organizations have fully scaled agentic AI deployment

Yet 93% of leaders believe scaling agents in the next 12 months provides a competitive edge. Everyone sees the prize; almost nobody has operationalized it.

**CAPGEMINI RESEARCH INSTITUTE, JULY
2025**

40%+

Of agentic AI projects will be cancelled by end of 2027

Escalating costs, unclear business value, inadequate risk controls: the reasons read like a governance checklist.

GARTNER, JUNE 2025

The diagnostic

Three questions your team should be able to answer, *but usually can't today*

If you added a new agent tomorrow, would it know everything your existing agents know?

If no Your memory is per-agent, and nothing you deploy compounds. You're buying tools, not building capability.

Can you name which of your AI investments moved a deal stage last quarter?

If no You have observability of cost, not outcome. You can't manage what you can't attribute, and you won't defend the budget next cycle.

Who approved the last thing an agent wrote to your CRM?

If nobody knows You don't have governance, you have exposure. This is the question the CFO asks at renewal.

Three no's is the normal answer today. That's the point: **it isn't a maturity failure, it's a missing layer.**

The through-line

Every theme points at the same four things *underneath the agents*

Grounding

In what's actually true about this account, this quarter, this product.

Capability

Reusable skills any agent can invoke, not rebuilt per tool.

Governance

On every action: who can act, on what, with what approval and audit trail.

Feedback

From every outcome, so the system learns instead of repeating.

This isn't a maturity failure. The models work. The stack around them doesn't, and it's the same missing stack every time. **Those four aren't features of any agent. They're the layer the agents run on.**

Where this points

wysdym is the *operating layer* for agentic go-to-market.

It grounds every agent in your GTM truth, runs the motion under your governance, and gets sharper with every deal you close. wysdym detects, directs, governs, and measures. It is the layer your agents and your team run on, not another agent. Every agent inherits the same grounding, acts within the limits you set, and gets smarter from every outcome.

Where every action compounds into a revenue advantage.