

52 agentic AI statistics *every GTM leader* should know in 2026

Agentic AI crossed from demo to deployment. These 52 statistics show where the market really stands. We've organized them around the five pillars that separate production agents from prototypes: **learning, skills, connectivity, governance, and observability**.

RC

Rob Catalano · Co-Founder, wysdym

Last updated: September 1, 2026

Every stat links to its original source

Live, always-current version with clickable sources: <https://gtmstats.wysdym.ai>

\$2.59T

Worldwide AI spending forecast for 2026 — up 47% in a single year. (*Gartner*)

53%

Of organizations are deploying AI agents — up from 12% in 2024 — and multi-agent orchestration doubled last quarter. (*KPMG*)

14%

Of teams ship agents with full security and IT approval — while 81% are already past planning. The five pillars below decide which side of this stat you land on. (*Gravitee*)

Every team is buying agents. No one is building the platform underneath.

MAJOR TAKEAWAYS

If you only read four things



Deployment is mainstream. Depth is rare.

53% of organizations now deploy AI agents and AI took over 70% of Q2 2026 venture funding — yet 95% of agents never reach production, and just 16% of enterprise deployments qualify as true agents.



The failure mode is knowledge, not models.

Agent task capability nearly quadrupled in a year — but 95% of GenAI pilots still return zero. MIT's diagnosis: systems that don't learn, don't retain context, and don't improve with use.



Governance became the deciding vote.

Human validation of agent outputs nearly tripled in a year (22% → 63%), yet only 14% of teams ship agents with full security and IT approval — and 88% of organizations report confirmed or suspected agent-related incidents.



Visibility is the trust currency.

89% of agent teams run some observability — 94% of those in production — but roughly 3 in 10 still don't evaluate their agents at all, while documented AI incidents rose 55% in 2025.

WHAT'S IN THIS REPORT

- 01 Market size & momentum
- 02 Adoption & implementation
- 03 The five pillars framework
- 04 Pillar 1 — Learning & knowledge
- 05 Pillar 2 — Skills & capabilities
- 06 Pillar 3 — Connectivity
- 07 Pillar 4 — Governance & trust
- 08 Pillar 5 — Observability
- 09 ROI, outcomes & what's next
- 10 What this means for GTM teams
- 11 FAQ
- 12 Methodology & sources

SECTION 01

Market size & momentum

The agentic AI market stopped being a forecast and started being a budget line. Here's how big it is, how fast it's growing, and where the money is going.

STAT 01

\$2.59T

Worldwide AI spending will hit \$2.59 trillion in 2026 — up 47% in a single year

Gartner's May 2026 forecast has AI software alone growing from \$282.9B in 2025 to \$453.2B in 2026. The budget debate is over; the execution debate is just starting.

[Gartner, May 2026](#)

STAT 02

\$1.3T

AI spending grows 31.9% a year through 2029 — reaching \$1.3 trillion

IDC pins the growth squarely on agentic AI — applications plus the systems needed to manage "agentic fleets" — exceeding 26% of all worldwide IT spending by 2029.

[IDC, August 2025](#)

STAT 03

\$581.7B

Global corporate AI investment reached \$581.7 billion in 2025 — up 130% year-over-year

Private AI investment hit \$344.7B (+127.5%), and US private investment (\$285.9B) ran 23x China's (\$12.4B), per Stanford's 2026 AI Index.

[Stanford HAI, 2026 AI Index, April 2026](#)

STAT 04

\$510B

Global venture funding hit a record \$510 billion in H1 2026 — with AI taking over 70% of Q2 capital

H1 2026 topped all of 2025 combined, and OpenAI plus Anthropic alone raised \$217 billion — 43% of all H1 startup funding. Investors aren't hedging on the agent thesis — they're concentrating on it.

[Crunchbase News, July 2026](#)

STAT 05

\$37B

Enterprises spent \$37 billion on generative AI in 2025 — 3.2x the year before

The largest share, \$19B, went to the application layer — the agents and AI products teams actually use — per Menlo Ventures' enterprise survey and market model.

[Menlo Ventures, December 2025](#)

STAT 06

\$450B

Agentic AI could drive ~30% of enterprise application software revenue by 2035 — over \$450 billion

That's Gartner's best-case scenario, up from 2% in 2025. Even the conservative read makes agents the defining software category of the next decade.

[Gartner, August 2025](#)

SECTION 02

Adoption & implementation

Nearly every enterprise is experimenting with agents. Far fewer have them in production — and the gap between piloting and scaling is where most of the story lives.

STAT 07

53%

53% of organizations are deploying AI agents — and multi-agent orchestration doubled in a single quarter

KPMG's Q2 2026 Pulse: agent deployment has held above half (up from 12% in 2024), while organizations orchestrating multiple agents across workflows jumped from 9% to 18% in one quarter. Leaders plan an average \$202M in AI investment over the next 12 months.

[KPMG AI Quarterly Pulse, June 2026](#)

STAT 08

62%

62% of organizations are experimenting with AI agents — but only 23% are scaling them anywhere

And in any single business function, no more than 10% are scaling agents. The gap between "trying agents" and "running on agents" is the defining stat of 2026.

[McKinsey, The State of AI, November 2025](#)

STAT 09

95%

95% of AI agents never make it to production

Capgemini's follow-up to its 2025 agentic AI research puts a hard number on the pilot-to-production gap. The difference between the 95% and the 5% is rarely the model — it's the operating layer underneath it.

[Capgemini Research Institute, April 2026](#)

STAT 10

40%

40% of enterprise applications will ship with task-specific AI agents by the end of 2026

Up from less than 5% in 2025. Agents are becoming a feature of software you already own — which raises the bar for the platform underneath them: the shared graph, skills, and governance every one of those agents has to run on.

[Gartner, August 2025](#)

STAT 11

2%

Only 2% of organizations have fully scaled agentic AI deployment

Yet 93% of leaders believe scaling agents in the next 12 months will provide a competitive edge. Everyone sees the prize; almost nobody has operationalized it.

[Capgemini Research Institute, July 2025](#)

STAT 12

16%

Just 16% of enterprise AI deployments qualify as true agents

Menlo's bar: systems where the model plans, acts, observes feedback, and adapts. Most of what's labeled "agent" is still a fixed-sequence workflow — the buyer-side version of what Gartner calls "agent washing."

[Menlo Ventures, December 2025](#)

THE FRAMEWORK

The five pillars of *production-ready* agentic AI

Most statistics roundups organize by market size and hype. We organized this one around the five questions that decide whether an agent survives contact with real buyers — the same five pillars we're building wysdym around. The data that follows makes the case for each one.

PILLAR 01

Learning

Agents are only as good as the knowledge underneath them — and whether it compounds.

PILLAR 02

Skills

What agents can actually do — typed capabilities, not vague promises.

PILLAR 03

Connectivity

Agents that can't reach your stack can't do real work.

PILLAR 04

Governance

Autonomy without approvals is just risk with better marketing.

PILLAR 05

Observability

If you can't see what your agent did, you can't trust what it does next.

PILLAR 01

Learning & knowledge

An agent without your knowledge is a generic model with your logo on it. The data is blunt: fragmented, stale knowledge is the single biggest thing standing between GTM teams and agents that actually work.

STAT 13

95%

95% of enterprise GenAI pilots deliver zero return — and the core barrier is learning

MIT Project NANDA's report puts it plainly: most systems "do not retain feedback, adapt to context, or improve over time." 66% of executives want AI that learns from feedback; 63% demand tools that retain context.

[MIT Project NANDA, The GenAI Divide, 2025](#)

STAT 14

60%

Organizations will abandon 60% of AI projects unsupported by AI-ready data through 2026

63% of data leaders either don't have — or aren't sure they have — the right data management practices for AI. The model isn't the bottleneck. The knowledge is.

[Gartner, February 2025](#)

STAT 15

60%

Sales reps spend 60% of their time on non-selling tasks

Hunting for the right deck, re-entering CRM notes, chasing internal approvals — most of a rep's week goes to everything except customers. The counterweight: 85% of reps with agents say AI frees them to focus on higher-value work.

[Salesforce Sales Statistics, February 2026](#)

STAT 16

84%

84% of data leaders say their data strategy needs a complete overhaul before their AI ambitions can succeed

70% believe their most valuable business insights sit inside the sliver of data that's inaccessible — siloed or unusable — and 89% of leaders with AI in production have already seen inaccurate or misleading outputs. That's the exact gap a structured knowledge layer exists to close.

[Salesforce, State of Data & Analytics, November 2025](#)

STAT 17

47%

47% of digital workers struggle to find the information they need to do their jobs

The average desk worker now juggles 11 applications, up from 6 in 2019. Gartner's survey is from 2023 — and every fragmentation trend since has pointed the wrong way.

[Gartner Digital Worker Survey, 2023](#)

STAT 18

3.4x

Grounding an LLM in a knowledge graph lifted answer accuracy from 16% to 54%

The data.world AI Lab benchmark asked GPT-4 enterprise data questions: direct queries scored 16% accuracy; the same questions over a knowledge graph scored 54%. Structure isn't a nice-to-have — it's the accuracy mechanism.

[Sequeda, Allemang & Jacob, arXiv, 2023](#)

STAT 19

22–94%

Hallucination rates across 26 top models ranged from 22% to 94% on a new accuracy benchmark

Stanford's 2026 AI Index shows even frontier models fail unpredictably when they're not grounded — one model's measured accuracy fell from over 90% to 14.4% under the benchmark's conditions.

[Stanford HAI, 2026 AI Index, April 2026](#)

STAT 20

53%

53% of B2B marketing leaders say fewer than half their sales enablement materials get used

Content that doesn't propagate to the moment of use is budget burned. Forrester's data is from 2023; nothing in the 2026 surveys suggests it's improved.

[Forrester, December 2023](#)

THE WYSDYM LENS

This is why wysdym is built around wysdymCortex — a typed knowledge graph of your business, 39 GTM domains, self-learning — rather than a flat document store. Every agent on the operating layer reads from the same graph and writes what it learns back, so the next agent starts smarter than the last one. Knowledge in folders decays; a graph graded by deal outcomes compounds.

PILLAR 02

Skills & capabilities

"Agent" covers everything from a glorified FAQ to software that closes tickets end-to-end. The capability data shows what agents genuinely do well today — and where autonomy still breaks.

STAT 21

77.3%

AI agents' success rate on real-world terminal tasks jumped from 20% to 77.3% in a year

Stanford's 2026 AI Index also has agents solving cybersecurity problems 93% of the time, up from 15% in 2024. Capability is compounding fast — the constraint is moving elsewhere.

[Stanford HAI, 2026 AI Index, April 2026](#)

STAT 22

30%

The best AI agent completed only 30% of real workplace tasks fully autonomously

CMU's TheAgentCompany benchmark simulates an actual software company. Agents handled simpler tasks well; long-horizon, multi-step work still defeated them. Autonomy has a ceiling — design for it.

[Carnegie Mellon et al., TheAgentCompany, NeurIPS 2025](#)

STAT 23

54%

54% of sellers have already used AI agents — nearly 9 in 10 plan to by 2027

Sellers expect agents to cut prospect research time by 34% and email drafting by 36%, per Salesforce's survey of 4,050 sales professionals.

[Salesforce, State of Sales, February 2026](#)

STAT 24

1.7x

Top-performing sales teams are 1.7x more likely to use agents for prospecting

And 94% of sales leaders with agents call them critical to meeting business demands. Agents are showing up first where revenue is measured.

[Salesforce, State of Sales, February 2026](#)

STAT 25

4.8 hrs

AI saves sellers 4.8 hours a week — but 72% of sales organizations fail to reinvest the time in high-value work

The teams that do reinvest it are 2.2x more likely to exceed customer-growth goals and 3.1x more likely to exceed lead-to-opportunity conversion goals, per Gartner's 2026 survey of 210 sales leaders. The skill isn't saving time; it's redeploying it.

[Gartner CSO & Sales Leader Conference, May 2026](#)

STAT 26

83%

83% of marketers say customers expect two-way conversations — but 69% can't respond promptly

And 81% would trust AI to respond to customers to help them scale. The conversation gap is now the most quantified problem in marketing.

[Salesforce, State of Marketing, February 2026](#)

THE WYSDYM LENS

Capabilities should be typed, testable, and shared — not improvised per agent. On wysdym, skills are the verbs: a library of typed GTM capabilities any agent can invoke over MCP, and Plays chain them into routines where every action ties to a measurable deal outcome. That's how agent-vs-headcount math stops being a guess.

PILLAR 03

Connectivity & interoperability

An agent that can't reach your CRM, your content, and your calendar is a demo. 2025–2026 is when the industry converged on open protocols for connecting agents to the systems where work actually happens.

STAT 27

10,000+

The Model Context Protocol passed 10,000+ public servers and 97M+ monthly SDK downloads in its first year

MCP — adopted by ChatGPT, Gemini, Microsoft Copilot, Cursor, and VS Code — was donated to the Linux Foundation's new Agentic AI Foundation in December 2025. Agent-to-tool connectivity now has a standard.

[Anthropic, December 2025](#)

STAT 28

150+

150+ organizations now back the Agent2Agent (A2A) interoperability protocol — triple its launch roster

Launched by Google in April 2025 with 50+ supporters and now governed by the Linux Foundation, A2A reached production-ready status in under a year — SDKs in five languages, backing from AWS, Cisco, Google, IBM, Microsoft, Salesforce, SAP, and ServiceNow. Agents are becoming a network, not a feature.

[The Linux Foundation, April 2026](#)

STAT 29

5x+

AI inference costs per agentic workflow will rise more than fivefold through 2028

Gartner calls it the "inference paradox": per-token prices keep falling, but multistep agents consume more — and often more expensive — tokens, so total cost climbs anyway. Wiring the same integrations and retrieval separately into every agent is exactly the cost curve a shared connection layer flattens.

[Gartner, August 2026](#)

STAT 30

51%

51% of sales leaders say tech silos delay or limit their AI initiatives

Connectivity isn't plumbing — it decides whether your agent can see the systems your revenue actually runs through.

[Salesforce, State of Sales](#)

THE WYSDYM LENS

This is why wysdym is built on MCP — the protocol that went from roughly 100K to 97M monthly installs in 16 months (stat 27). Bring any agent — Claude, OpenAI, LangGraph, or your own — and one operating layer connects it to the systems your revenue runs through. Connections don't just feed answers; they feed your graph. Vendor-agnostic by design.

THE BET

The operating layer, *not the agent*

You bring the agent — Claude, OpenAI, LangGraph, or your own. wysdym brings what's underneath: shared graph memory, typed skills, connections into your stack, approval queues, and a feedback loop that learns from deal outcomes. The more agents you run, the smarter every one gets.

PILLAR 04

Governance & trust

The fastest way to kill an agent program is an agent that acts without guardrails. The trust data explains why buyers hesitate — and why the winners gate consequential actions instead of hoping for the best.

STAT 31

14%

Only 14% of teams ship AI agents with full security and IT approval

81% of teams have agents past the planning phase — in active testing or production — but just 14% go live with full sign-off from security and IT, and 88% of organizations report confirmed or suspected agent-related incidents. Adoption is outpacing control, and governance is the bottleneck. Survey of 900+ executives and practitioners.

[Gravitee, State of AI Agent Security 2026](#)

STAT 32

65%

65% of enterprises have seen an AI agent act outside its intended scope

And 29% took measurable business impact from it. The confidence gap is the tell: 94% of leaders trust their agents aren't over-provisioned, yet only a third actually enforce least-privilege access — and nearly half couldn't produce a complete 30-day audit trail. Trust without enforcement isn't governance.

[EMA for Cequence Security, "Agents Without Guardrails," August 2026](#)

STAT 33

40%

By 2027, 40% of enterprises will demote or decommission AI agents over governance gaps found only after production incidents

Gartner's May 2026 warning is specific: treating agent governance as binary — locked down or fully trusted — is the root cause of failure. Graduated autonomy wins.

[Gartner, May 2026](#)

STAT 34

63%

63% of leaders now require human validation of AI agent outputs — nearly triple a year ago

Human-in-the-loop went from afterthought to default in four quarters (22% in Q1 2025 → 63% in Q1 2026), and 91% say data security, privacy, and risk concerns now shape their AI strategy.

[KPMG AI Quarterly Pulse, March 2026](#)

STAT 35

1 in 4

Only one-quarter of leaders completely trust their AI systems — yet two-thirds have given AI write access to core systems

And more than four in five tech leaders expect autonomous agents to make decisions with material business impact within a year, per Kyndryl's survey of 1,100 business and technology leaders. Trust is trailing autonomy — the market is voting for governed autonomy.

[Kyndryl People Readiness Report, June 2026](#)

STAT 36

21%

Nearly 75% of companies plan to deploy agentic AI within two years — only 21% have mature agent governance

Deloitte's 2026 State of AI finds the most successful companies start with lower-risk use cases, build governance capability, and scale deliberately. Governance maturity is the adoption rate-limiter.

[Deloitte, State of AI in the Enterprise, January 2026](#)

STAT 37

1 in 5

1 in 5 organizations suffered a breach tied to shadow AI — adding \$670K to average breach costs

And 63% of breached organizations had no AI governance policy at all. Ungoverned agents aren't just a quality risk; they're a security line item.

[IBM, Cost of a Data Breach, July 2025](#)

STAT 38

**Aug
2026**

From August 2, 2026, EU rules require companies to disclose when users are interacting with AI

The EU AI Act's transparency obligations proceed on schedule even as high-risk deadlines move to December 2027 under the May 2026 omnibus agreement (pending formal adoption). Most consumers want this anyway — nearly 75% want to know when they're talking to an AI agent (Salesforce).

[Gibson Dunn / EU AI Act tracker, May 2026](#)

THE WYSDYM LENS

Our rule is simple: every agent action visible, the consequential ones gated. On wysdym, agents write through approval queues — a human signs off before anything touches your CRM. Approvals aren't friction; they're what makes autonomy deployable. Teams shouldn't have to choose between an agent that does nothing and an agent they can't control.

PILLAR 05

Observability

Would you let a new hire touch every deal in your pipeline without ever reviewing the work? That's what an unobserved agent is — and once you're running several, the surface area only grows. The market is waking up to this, fast.

STAT 39

89%

89% of organizations run some observability on their agents — but only 62% can trace individual steps and tool calls

Among teams already in production it's 94% and 71.5%. The teams that ship treat visibility as non-negotiable, not nice-to-have.

[LangChain, State of Agent Engineering, December 2025](#)

STAT 40

3 in 10

Roughly 3 in 10 agent teams aren't evaluating their agents at all

Only 52.4% run offline evals before deployment, and just 37.3% evaluate live traffic. Most agent failures aren't mysterious — they're unmeasured.

[LangChain, State of Agent Engineering, December 2025](#)

STAT 41

15→50%

LLM observability investments will grow from 15% of GenAI deployments to 50% by 2028

Gartner ties the jump to explainable AI becoming a requirement for secure deployment. A separate May 2026 Gartner prediction: 40% of AI-deploying organizations will run dedicated AI observability tools by 2028.

[Gartner, March 2026](#)

STAT 42

+55%

Documented AI incidents hit a record 362 in 2025 — up 55% in one year

Stanford's AI Index tracks them through the AI Incident Database. More agents, more autonomy, more surface area — visibility is how you stay off this list.

[Stanford HAI, 2026 AI Index, April 2026](#)

STAT 43

54/yr

Enterprises average 54 AI agent incidents a year that require human correction

37% of those incidents caused data exposure, 33% triggered cascading failures, and 17% created compliance issues — while organizations with embedded, automated controls deploy 16x more agents than those governing manually.

[IBM Institute for Business Value, June 2026](#)

STAT 44

5%

About 5% of AI model requests fail in production

Datadog's telemetry from 1,000+ companies found ~60% of those failures come from provider rate limits — and 70%+ of organizations now run three or more models. Real systems fail in boring ways. You need to see it happen.

[Datadog, State of AI Engineering, April 2026](#)

STAT 45

78%

78% of executives lack confidence they could pass an independent AI governance audit within 90 days

Nearly 3 in 4 organizations are giving agents access to their data and processes — but just 20% have a tested AI incident response plan. Auditability is the new uptime.

[Grant Thornton, 2026 AI Impact Survey](#)

THE WYSDYM LENS

Observability is how agents earn trust — and how they get smarter. wysdym is designed to trace every agent action and attribute it to deal outcomes: Findings spot what's drifting before deals stall, and everything observed flows back into the graph. Visibility doesn't just catch problems; it compounds learning.

SECTION 09

ROI, outcomes & what's next

The honest version: returns are real but unevenly distributed. The teams seeing ROI aren't the ones with the most pilots — they're the ones whose agents learn, connect, and stay governed.

STAT 46

39%

88% of organizations use AI — but only 39% see any EBIT impact from it

And most of those attribute less than 5% of EBIT to AI. The bright spot: revenue gains are reported most often in marketing and sales use cases.

[McKinsey, The State of AI, November 2025](#)

STAT 47

7%

84% of finance organizations have implemented or plan to implement AI — but only 7% report high impact

Gartner's survey of 183 CFOs is the CFO-side version of the depth problem: near-universal adoption, single-digit transformation. The gap between the 84% and the 7% is an operating-model gap — compare stat 11 (2% fully scaled) and stat 12 (16% true agents).

[Gartner, June 2026](#)

STAT 48

\$4.4M

99% of large enterprises reported financial losses from AI-related risks — averaging \$4.4 million each

64% lost over \$1 million. EY's survey of 975 C-suite leaders is the clearest price tag yet on ungoverned, unobserved AI.

[EY Responsible AI Pulse, October 2025](#)

STAT 49

67%

67% of B2B buyers prefer a rep-free buying experience

And 45% used AI during a recent purchase. Your buyers brought agents to the table first — the question is whether the agents on your side are grounded, governed, and learning, or improvising in isolation.

[Gartner, March 2026](#)

STAT 50

7x

Contacting a web lead within an hour makes you ~7x more likely to qualify them

The classic HBR study (2011) is still the benchmark: 60x better than waiting 24+ hours — yet average response time was 42 hours, and 23% of companies never responded at all. Speed-to-lead remains the oldest argument for always-on agents — and for grounding them so the first touch is right, not just fast.

[Harvard Business Review, 2011](#)

STAT 51

\$262B

AI and agents drove 20% of global online orders — \$262 billion — during the 2025 holiday season

Salesforce's commerce data is the largest real-world signal yet that buyers will transact through agentic experiences, not just research through them.

[Salesforce, February 2026](#)

STAT 52

60%

By 2028, 60% of brands will use agentic AI for one-to-one buyer interactions

Gartner calls it the end of channel-based marketing: instead of campaigns pushed through channels, agents handle individual conversations. GTM's next operating model is conversational.

[Gartner, January 2026](#)

What this means for GTM teams

Read all 52 together and one pattern keeps surfacing: capability stopped being the constraint. The operating model is.

Agent task performance nearly quadrupled in a year (stat 21) — yet 95% of pilots return zero (stat 13), only 14% of teams ship agents with full security approval (stat 31), and only a quarter of leaders completely trust their AI systems (stat 35). The teams winning aren't the ones with better models. Everyone has the same models.

What separates them maps cleanly to five pillars. They ground agents in structured, learning knowledge instead of flat documents (stats 13, 14, 18). They scope skills to what agents verifiably do well (stat 22). They connect agents to the systems where work actually happens (stats 27–30). They gate consequential actions with human judgment (stat 34). And they watch everything (stats 39–40).

For GTM teams, the urgency is sharper. 67% of your buyers prefer to buy without a rep (stat 49). They expect two-way conversations your team can't staff (stat 26). They reward whoever answers first (stat 50). Agents will do more and more of this work — so the question the data keeps asking isn't which agent you'll run. It's what your agents run on.

That's the bet behind wysdym: the operating layer, not the agent. Every team is buying agents; no one is building the platform underneath. You bring the agent — Claude, OpenAI, LangGraph, or your own. wysdym is the platform underneath: shared graph memory, typed skills, approval-gated writes, and a feedback loop that ties every agent action to deal outcomes — so the next agent on your stack starts smarter than the last one.

Frequently asked questions

What is agentic AI?

AI systems that don't just generate answers — they plan and take actions toward goals: retrieving knowledge, using tools, and executing multi-step work with varying levels of autonomy. Menlo Ventures' working definition is a useful bar: a true agent plans, acts, observes feedback, and adapts. By that standard, only 16% of enterprise AI deployments qualify today (stat 12).

How is an AI agent different from a chatbot?

A scripted chatbot follows a decision tree and breaks the moment a buyer goes off-script. An agent reasons over knowledge and takes action — and a knowledge-grounded agent answers from your company's actual content, positioning, and intelligence, so the answers are yours rather than a generic model's. The difference is measurable: grounding answers in structured knowledge took accuracy from 16% to 54% in benchmark testing (stat 18).

How big is the agentic AI market in 2026?

Worldwide AI spending reaches \$2.59 trillion in 2026 (Gartner), and 53% of organizations are deploying agents (KPMG). On the agentic-specific slice: Gartner projects agentic AI could drive over \$450 billion in enterprise application software revenue by 2035, and Capgemini estimates up to \$450 billion in economic value from agents by 2028.

Why do most agentic AI projects fail?

Three patterns dominate the research. Knowledge: 95% of pilots return zero because systems don't learn or retain context (MIT Project NANDA), and 60% of AI projects get abandoned without AI-ready data (Gartner). Governance: only 14% of teams ship agents with full security and IT approval, and 88% of organizations report agent-related incidents (Gravitee). Observability: roughly 3 in 10 teams don't evaluate their agents at all (LangChain). Failed agent programs almost never fail on model quality.

What should GTM leaders look for in the platform their agents run on?

Ask five questions. **Learning:** do agents run on your structured knowledge, and does it get smarter with every outcome? **Skills:** are capabilities typed and shared across agents — and honest about what they don't do? **Connectivity:** does one layer plug every agent into your CRM, content, and conversation tools? **Governance:** can you gate consequential writes behind human approval? **Observability:** can you trace every action and tie it back to deal outcomes?

What does wysdym do?

wysdym is the operating layer for agentic GTM — the layer underneath your agents, not another agent. You bring any agent (Claude, OpenAI, LangGraph, or custom), and wysdym supplies what every agent needs to work: shared memory through a typed knowledge graph of your business (wysdymCortex), an open skill library any agent invokes over MCP (wysdymSkills), your stack — Salesforce, HubSpot, Gong, Slack, Drive, Notion — wired once, not per agent (wysdymConnect), one governed door with per-agent permissions, human-in-the-loop, and receipts (wysdymGateway), and outcome attribution to deal stages, with Findings that flag drift with the fix attached (wysdymObserve). The more agents you run on it, the smarter every one gets.

Is wysdym the same company as Wysdom.AI?

No. They are two unrelated companies with similar names. **wysdym** — spelled w-y-s-d-y-m — is the operating layer for agentic GTM, founded in 2025 in Toronto by Rob Catalano and Nadir Ebrahim, at wysdym.ai. **Wysdom.AI** — spelled w-y-s-d-o-m — is a conversational-AI and bot-analytics company acquired by Calabrio in January 2024. There is no affiliation, ownership, or shared history between them. If a search engine or AI assistant describes wysdym as a chatbot-analytics or contact-centre vendor, it has merged the two companies.

METHODOLOGY

How we compiled this

All 52 statistics on this page link to their original sources — analyst press releases, primary research reports, vendor surveys with disclosed methodology, peer-reviewed papers, and official announcements. We verified each figure against the original publication in September 2026, and where a widely-cited number is older than 2025 (like the HBR speed-to-lead study), we say so instead of pretending it's new.

None of the numbers on this page are wysdym's own. We'd rather show you the industry's real data than dress up our opinions as research. Where we editorialize, it's clearly labeled as "the wysdym lens."

Found an error or a fresher source? [Tell us](#) — we update this page as the research evolves.

PUT THE PILLARS TO WORK

Agents alone won't grow revenue.

Compounding intelligence will.

Any vendor can sell you an agent. None of them make your other agents better, act on your terms, or tell you which ones are working. That's where wysdym comes in — [the operating layer underneath](#), where every action compounds into a revenue advantage.



The operating layer for agentic GTM. All third-party statistics remain the property of their respective publishers; each is linked to its original source. Compiled July 2026
· Updated September 2026.

Go deeper on the wysdym blog: [Why We're Building the Operating Layer, Not Another Agent](#) · [95% of Enterprise AI Pilots Return Zero](#) · [The 40% Problem](#) · [Your AI Agents Don't Share a Brain](#)

On wysdym.ai: [The platform: five pillars, one operating layer](#) · [Go-to-market intelligence](#) · [How to evaluate GTM agent platforms](#)

Field notes: [What ~100 GTM Leaders Told Us About AI Agents — and the Data Behind It](#)